

Speech and Language Processing

Chapter 6: Hidden Markov Models (HMM)

Ines Rehbein

NCLT, Dublin City University



Outline


- 1 Hidden Markov Models
 - Introduction
 - Evaluation: the Forward Algorithm
 - Decoding: the Viterbi Algorithm
 - Training: the Forward-Backward Algorithm

Overview

- HMMs and MEMs are probabilistic sequence classifiers:
given a sequence of units they compute a probability distribution over possible labels and choose the best label sequence

Background: Extensions of Finite Automata

- The **HMM** is an extensions of a finite automaton
- **Finite automata**: set of states, set of transitions between states taken based on the input



- **Weighted finite-state automata**: each arc is associated with a probability
- **Markov chain**:
 - weighted automaton in which input uniquely determines the transitions
 - can only deal with unambiguous input data

Hidden Markov Model

- HMM can deal with “hidden” data (data which cannot be observed directly)
- HMM consists of:
 - $Q = q_1 q_2 \dots q_N$ a set of N **states**
 - $A = a_{11} a_{12} \dots a_{n1} \dots a_{nn}$ a **transition probability matrix**
 - $O = o_1 o_2 \dots o_T$ a sequence of **observations**
 - $B = b_i(o_t)$ a sequence of **observation likelihoods**
 - q_0, q_F a **start state** and a **final state**
- fully connected (ergodic) HMM vs. left-to-right (Bakis) HMM

First Order HMM

- 2 simplifying assumptions:

- ① **Markov Assumption:** the probability of a particular state is dependent only on the previous state

$$P(q_i | q_1 \dots q_{i-1}) = P(q_i | q_{i-1}) \quad (1)$$

- ② **Output Independence Assumption:** the probability of an output observation is dependent only on the state that produced the observation


$$P(o_i | q_1 \dots q_i, \dots, q_T, o_1, \dots, o_i, \dots, o_T) = P(o_i | q_i) \quad (2)$$

3 Main Problems

- ① **Evaluation:** Given an HMM $\lambda = (A, B)$ and an observation sequence O , determine the likelihood $P(O|\lambda)$
 - Which HMM most probably generated a given sequence O ?
 - Application: Speech recognition
- ② **Decoding:** Given an observation sequence O and an HMM $\lambda = (A, B)$, discover the best hidden state sequence Q
 - Application: Tagging
- ③ **Learning:** Given an observation sequence O and the set of states in the HMM, learn the HMM parameters A and B
 - Generate an HMM from a sequence of observations
 - Application: Learn the parameters of an HMM (Training)

Evaluation: the Forward Algorithm

- **Task:** We have a sequence of observations O telling us how many icecreams Jason Eisner ate in summer 2007.
⇒ Determine the probability of an icecream observation without knowing the hidden state sequence (weather conditions)



Evaluation: the Forward Algorithm

- **Simplify** task: compute the probability of an icecream observation O given a state sequence Q :

$$P(O|Q) = \prod_{i=1}^T P(o_i|q_i) \quad (3)$$

- Probability for icecream observation (3, 1, 3) with state sequence (hot hot cold):

$$P(3 \ 1 \ 3 | \text{hot hot cold}) = P(3|\text{hot}) \times P(1|\text{hot}) \times P(3|\text{cold}) \quad (4)$$

Evaluation: the Forward Algorithm

- **Simplify** task: compute the probability of an icecream observation O given a state sequence Q :

$$P(O|Q) = \prod_{i=1}^T P(o_i|q_i) \quad (3)$$

- Probability for icecream observation (3, 1, 3) with state sequence (hot hot cold):

$$P(3 \ 1 \ 3 | \text{hot hot cold}) = P(3|\text{hot}) \times P(1|\text{hot}) \times P(3|\text{cold}) \quad (4)$$

Evaluation: the Forward Algorithm (II)

- But: we don't know the hidden state sequence
⇒ compute probability of observation (3, 1, 3) by summing over all possible weather sequences, weighted by their probability
- Compute the joint probability of being in a particular weather sequence Q and generating a particular sequence O of ice-cream events:

$$P(O, Q) = P(O|Q) \times P(Q) = \prod_{i=1}^n P(o_i|q_i) \times \prod_{i=1}^n P(q_i|q_{i-1}) \quad (5)$$

$$P(3 \ 1 \ 3, \text{ hot hot cold}) =$$

$$P(3 \ 1 \ 3, \text{ cold cold cold}) + P(3 \ 1 \ 3, \text{ cold cold hot}) + \\ P(3 \ 1 \ 3, \text{ hot hot cold}) + \dots$$

Evaluation: the Forward Algorithm (II)

- But: we don't know the hidden state sequence
⇒ compute probability of observation (3, 1, 3) by summing over all possible weather sequences, weighted by their probability
- Compute the joint probability of being in a particular weather sequence Q and generating a particular sequence O of ice-cream events:

$$P(O, Q) = P(O|Q) \times P(Q) = \prod_{i=1}^n P(o_i|q_i) \times \prod_{i=1}^n P(q_i|q_{i-1}) \quad (5)$$

$$P(3 \ 1 \ 3, \text{ hot hot cold}) =$$

$$P(3 \ 1 \ 3, \text{ cold cold cold}) + P(3 \ 1 \ 3, \text{ cold cold hot}) + \\ P(3 \ 1 \ 3, \text{ hot hot cold}) + \dots$$

Evaluation: the Forward Algorithm (II)

- But: we don't know the hidden state sequence
⇒ compute probability of observation (3, 1, 3) by summing over all possible weather sequences, weighted by their probability
- Compute the joint probability of being in a particular weather sequence Q and generating a particular sequence O of ice-cream events:

$$P(O, Q) = P(O|Q) \times P(Q) = \prod_{i=1}^n P(o_i|q_i) \times \prod_{i=1}^n P(q_i|q_{i-1}) \quad (5)$$

$$P(3 \ 1 \ 3, \textit{hot hot cold}) =$$

$$P(3 \ 1 \ 3, \textit{cold cold cold}) + P(3 \ 1 \ 3, \textit{cold cold hot}) + \\ P(3 \ 1 \ 3, \textit{hot hot cold}) + \dots$$

Evaluation: the Forward Algorithm (III)

- For an HMM with N hidden states and T observations
 $\Rightarrow N^T$ possible hidden sequences
- For a real task N^T is too large to compute the observation likelihood for each hidden state
- Instead: Use **Forward Algorithm** (dynamic programming algorithm with $O(TN^2)$)

Evaluation: the Forward Algorithm (IV)

- Forward Algorithm computes the observation probability by **summing over the probabilities of all possible hidden state paths** that could generate the observation sequence
- Forward Algorithm efficiently stores the probabilities for each path in a **forward trellis**
- Each cell of the trellis $\alpha_t(j)$ represents the probability of being in state j after seeing the first t observations, given the automaton λ
- The value of each cell is computed by summing over the probabilities of every path that could lead to this cell:

$$\alpha_t(j) = P(o_1, o_2 \dots o_t, q_t = j | \lambda) \quad (6)$$

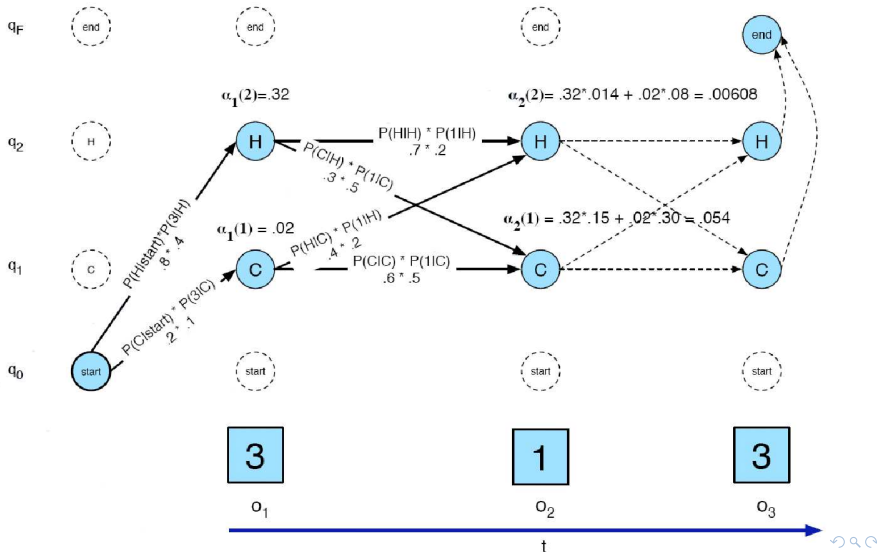
Evaluation: the Forward Algorithm (V)

- For a given state q_j at time t , the value $\alpha_t(j)$ is computed as:

$$\alpha_t(j) = \sum_{i=1}^N a_{t-1}(i) a_{ij} b_j(o_t) \quad (7)$$

- $\alpha_{t-1}(i)$ the **previous forward path probability** from the previous time step
- a_{ij} the **transition probability** from previous state q_i to current state q_j
- $b_j(o_t)$ the **state observation likelihood** of the observation symbol o_t given the current state j

Evaluation: the Forward Algorithm (VI)



Evaluation: the Forward Algorithm (VII)

function FORWARD(*observations* of len T , *state-graph* of len N) **returns** *forward-prob*

create a probability matrix $forward[N+2, T]$

for each state s **from** 1 **to** N **do** ; initialization step

$forward[s, 1] \leftarrow a_{0,s} * b_s(o_1)$

for each time step t **from** 2 **to** T **do** ; recursion step

for each state s **from** 1 **to** N **do**

$forward[s, t] \leftarrow \sum_{s'=1}^N forward[s', t-1] * a_{s',s} * b_s(o_t)$

$forward[q_F, T] \leftarrow \sum_{s=1}^N forward[s, T] * a_{s,q_F}$; termination step

return $forward[q_F, T]$

The Forward Algorithm - Summary

- **Task:** Find the probability of a sequence of observations O given an HMM
- We **reduce the complexity** of calculating the probability by creating a **trellis** and calculating **partial probabilities** for each cell in the trellis (the probability of getting to a particular state q at time t)
- The probability of the observation sequence O is computed **recursively** by calculating the partial probabilities at time $t = 1, 2, \dots, T$ and adding all α 's at $t = T$

Decoding: The Viterbi Algorithm

- **Task:** We have a sequence of icecream observations (3, 1, 3) and an HMM. Find the sequence of hidden states which most likely produced the sequence of observations (3, 1, 3)
- We can't use the Forward Algorithm because there is an **exponentially large** number of state sequences
- **Solution:** Use Viterbi algorithm (dynamic programming algorithm with $O(TN^2)$)

Decoding: The Viterbi Algorithm (II)

- **Idea:** process the observation sequence left-to-right, fill the trellis
- Each cell of the trellis $v_t(j)$ represents the probability that the HMM is in state j after seeing the first t observations and passing through the **most probable** state sequence q_0, q_1, \dots, q_{t-1} , given the automaton λ
- The value of $v_t(j)$ is computed by recursively taking the most probable path that could lead to this cell:

$$v_t(j) = \max_{q_0, q_1, \dots, q_{t-1}} P(q_0, q_1 \dots q_{t-1}, o_1, o_2 \dots o_t, q_t = j | \lambda) \quad (8)$$

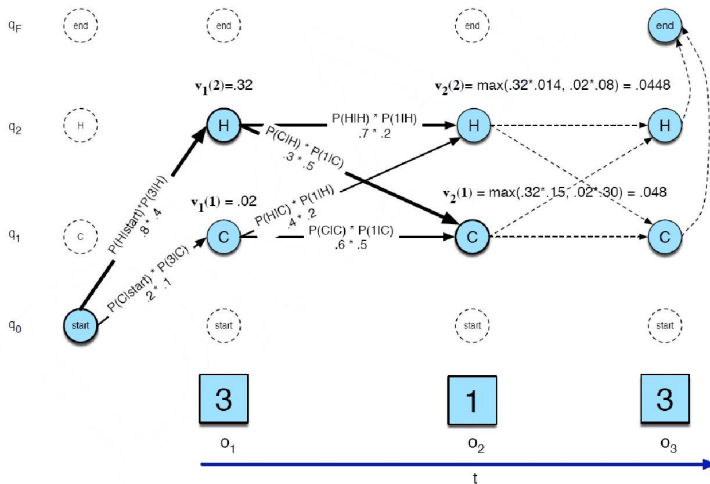
Decoding: The Viterbi Algorithm (III)

- For a given state q_j at time t , the value $v_t(j)$ is computed as:

$$v_t(j) = \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t) \quad (9)$$

- $v_{t-1}(i)$ the **previous Viterbi path probability** from the previous time step
- a_{ij} the **transition probability** from previous state q_i to current state q_j
- $b_j(o_t)$ the **state observation likelihood** of the observation symbol o_t given the current state j

Computing Likelihood: the Viterbi Algorithm (IV)



Computing Likelihood: the Viterbi Algorithm (VI)

- Find the **most probable** sequence of hidden states given a sequence of observed states
- Exploit **time invariance** of probabilities to reduce complexity by avoiding the necessity for examining every route through the trellis
- Keep a **backward pointer** for each state ($t > 1$), and store probability with each state (probability of having reached the state following the path indicated by the back pointers)
- When the algorithm reaches the states at time $t = T$, the probabilities for the final states are the probabilities of following the most probable route to that state

Computing Likelihood: the Viterbi Algorithm (V)

function VITERBI(*observations* of len T , *state-graph* of len N) **returns** *best-path*

create a path probability matrix $viterbi[N+2, T]$

for each state s **from** 1 **to** N **do** ;initialization step

$viterbi[s, 1] \leftarrow a_{0,s} * b_s(o_1)$

$backpointer[s, 1] \leftarrow 0$

for each time step t **from** 2 **to** T **do** ;recursion step

for each state s **from** 1 **to** N **do**

$viterbi[s, t] \leftarrow \max_{s'=1}^N viterbi[s', t-1] * a_{s',s} * b_s(o_t)$

$backpointer[s, t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s', t-1] * a_{s',s}$

$viterbi[q_F, T] \leftarrow \max_{s=1}^N viterbi[s, T] * a_{s,q_F}$; termination step

$backpointer[q_F, T] \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s, T] * a_{s,q_F}$; termination step

return the backtrace path by following backpointers to states back in time from $backpointer[q_F, T]$

The Viterbi Algorithm - Summary

- **Viterbi** is identical to the forward algorithm except that it takes the **max** over the previous path probabilities where the **forward** algorithm takes the **sum**
- Forward algorithm computes observation likelihood, Viterbi computes **most probable state sequence**
- Viterbi has **backpointers** to keep track of the path of hidden states that leads to each state
⇒ trace back the best path that leads to each state
(Viterbi **backtrace**)

Training: The Forward-Backward Algorithm (Baum-Welch Algorithm)

- **Task:** We have a sequence of icecream observations O and a set of hidden states H, C
⇒ Train the HMM and learn the parameters A (transition probabilities) and B (observation likelihood)
- Forward-Backward Algorithm: Overview
 - Make an **initial guess** at the parameters, then assess the approximation and try to reduce the error
 - Compute the **forward probability** of arriving at the state given the approximation and the **backward probability** of generating the final state of the model

Training: The Forward-Backward Algorithm (Baum-Welch Algorithm)

- **Task:** We have a sequence of icecream observations O and a set of hidden states H, C
⇒ Train the HMM and learn the parameters A (transition probabilities) and B (observation likelihood)
- Forward-Backward Algorithm: Overview
 - Make an **initial guess** at the parameters, then assess the approximation and try to reduce the error
 - Compute the **forward probability** of arriving at the state given the approximation and the **backward probability** of generating the final state of the model

Training: The Forward-Backward Algorithm (II)

- If we would know which hidden states produced the output, we could compute the **maximum likelihood estimate** of the transition probability a_{ij} :
count the number of times the transition was taken $C(i \rightarrow j)$, and normalise by the total count of all times we took any transition from state i :

$$a_{ij} = \frac{C(i \rightarrow j)}{\sum_{q \in Q} C(i \rightarrow q)} \quad (10)$$

- Can be directly computed for a Markov chain, but not for an HMM (we don't know which path of states was taken to get a particular input)

Training: The Forward-Backward Algorithm (II)

- If we would know which hidden states produced the output, we could compute the **maximum likelihood estimate** of the transition probability a_{ij} :
count the number of times the transition was taken $C(i \rightarrow j)$, and normalise by the total count of all times we took any transition from state i :

$$a_{ij} = \frac{C(i \rightarrow j)}{\sum_{q \in Q} C(i \rightarrow q)} \quad (10)$$

- Can be directly computed for a Markov chain, but not for an HMM (we don't know which path of states was taken to get a particular input)

Training: The Forward-Backward Algorithm (III)

- Forward-Backward Algorithm uses two intuitions to solve the problem:
 - 1 *iteratively* estimate the counts: start with an estimate for the transition and observation probabilities, then use these estimated probabilities to derive better and better probabilities
 - 2 estimate probabilities by computing the forward probability for an observation and dividing that probability mass among all the different paths that contributed to this forward probability

Training: The Forward-Backward Algorithm (V)

- Define **backward probability** β : probability of seeing the observations from time $t + 1$ to the end, given that we are in state j at time t , given the automaton λ

$$\beta_t(i) = P(o_{t+1}, o_{t+2} \dots o_T | q_t = i, \lambda) \quad (11)$$

- Sum over all successive values $\beta_{t+1}(j)$ weighted by their transition probabilities a_{ij} and observation probabilities $b_j(o_{t+1})$

- Initialisation:** $\beta_T(i) = a_{i,F}, 1 \leq i \leq N$ (12)

- Recursion:**
$$\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(o_{t+1}) \beta_{t+1}(j), 1 \leq i \leq N, 1 \leq t < T \quad (13)$$

- Termination:**
$$P(O|\lambda) = \alpha_T(q_F) = \beta_1(0) = \sum_{j=1}^N a_{0j} b_j(o_1) \beta_1(j) \quad (14)$$

Training: The Forward-Backward Algorithm (V)

- Define **backward probability** β : probability of seeing the observations from time $t + 1$ to the end, given that we are in state j at time t , given the automaton λ

$$\beta_t(i) = P(o_{t+1}, o_{t+2} \dots o_T | q_t = i, \lambda) \quad (11)$$

- Sum over all successive values $\beta_{t+1}(j)$ weighted by their transition probabilities a_{ij} and observation probabilities $b_j(o_{t+1})$

- Initialisation:** $\beta_T(i) = a_{i,F}, 1 \leq i \leq N$ (12)

- Recursion:**

$$\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(o_{t+1}) \beta_{t+1}(j), 1 \leq i \leq N, 1 \leq t < T \quad (13)$$

- Termination:**

$$P(O|\lambda) = \alpha_T(q_F) = \beta_1(0) = \sum_{j=1}^N a_{0j} b_j(o_1) \beta_1(j) \quad (14)$$

Training: The Forward-Backward Algorithm (V)

- **Forward** and **backward probability** help to compute the transition probability a_{ij} and observation probability $b_i(o_t)$ from an observation sequence, even though the actual path is hidden
- Estimate the probability a_{ij} of a particular transition between states i and j :

$$\hat{a}_{ij} = \frac{\text{expected number of transitions from state } i \text{ to state } j}{\text{expected number of transitions from state } i}$$

- How do we get the numerator?

If we had an estimate of the probability that a transition $i \rightarrow j$ was taken at time t for each $t \in T$, we could sum over all times t to estimate the total count:

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j | O, \lambda) \quad (15)$$

Training: The Forward-Backward Algorithm (V)

- **Forward** and **backward probability** help to compute the transition probability a_{ij} and observation probability $b_i(o_t)$ from an observation sequence, even though the actual path is hidden
- Estimate the probability a_{ij} of a particular transition between states i and j :

$$\hat{a}_{ij} = \frac{\text{expected number of transitions from state } i \text{ to state } j}{\text{expected number of transitions from state } i}$$

- How do we get the numerator?

If we had an estimate of the probability that a transition $i \rightarrow j$ was taken at time t for each $t \in T$, we could sum over all times t to estimate the total count:

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j | O, \lambda) \quad (15)$$

Training: The Forward-Backward Algorithm (V)

- **Forward** and **backward probability** help to compute the transition probability a_{ij} and observation probability $b_i(o_t)$ from an observation sequence, even though the actual path is hidden
- Estimate the probability a_{ij} of a particular transition between states i and j :

$$\hat{a}_{ij} = \frac{\text{expected number of transitions from state } i \text{ to state } j}{\text{expected number of transitions from state } i}$$

- How do we get the numerator?

If we had an estimate of the probability that a transition $i \rightarrow j$ was taken at time t for each $t \in T$, we could sum over all times t to estimate the total count:

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j | O, \lambda) \quad (15)$$

Training: The Forward-Backward Algorithm (V)

- **Forward** and **backward probability** help to compute the transition probability a_{ij} and observation probability $b_i(o_t)$ from an observation sequence, even though the actual path is hidden
- Estimate the probability a_{ij} of a particular transition between states i and j :

$$\hat{a}_{ij} = \frac{\text{expected number of transitions from state } i \text{ to state } j}{\text{expected number of transitions from state } i}$$

- How do we get the numerator?

If we had an estimate of the probability that a transition $i \rightarrow j$ was taken at time t for each $t \in T$, we could sum over all times t to estimate the total count:

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j | O, \lambda) \quad (15)$$

Training: The Forward-Backward Algorithm (VI)

- Define ξ_t as the probability of being in state i at time t and state j at time $t + 1$, given the observation sequence O and the HMM λ
- First compute something close to ξ , but including the probability of the observation:

$$\text{not - quite - } \xi_t(i, j) = P(q_t = i, q_{t+1} = j, O | \lambda) \quad (16)$$

- Use forward probability, transition probability, observation likelihood and backward probability:

$$\text{not - quite - } \xi_t(i, j) = \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j) \quad (17)$$

- To get ξ from not-quite- ξ , we have to divide by $P(O | \lambda)$, following the laws of probability:

$$P(X | Y, Z) = \frac{P(X, Y | Z)}{P(Y | Z)} \quad (18)$$

Training: The Forward-Backward Algorithm (VI)

- Define ξ_t as the probability of being in state i at time t and state j at time $t + 1$, given the observation sequence O and the HMM λ
- First compute something close to ξ , but including the probability of the observation:

$$\text{not - quite - } \xi_t(i, j) = P(q_t = i, q_{t+1} = j, O | \lambda) \quad (16)$$

- Use forward probability, transition probability, observation likelihood and backward probability:

$$\text{not - quite - } \xi_t(i, j) = \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j) \quad (17)$$

- To get ξ from not-quite- ξ , we have to divide by $P(O | \lambda)$, following the laws of probability:

$$P(X | Y, Z) = \frac{P(X, Y | Z)}{P(Y | Z)} \quad (18)$$

Training: The Forward-Backward Algorithm (VI)

- Define ξ_t as the probability of being in state i at time t and state j at time $t + 1$, given the observation sequence O and the HMM λ
- First compute something close to ξ , but including the probability of the observation:

$$\text{not - quite - } \xi_t(i, j) = P(q_t = i, q_{t+1} = j, O | \lambda) \quad (16)$$

- Use forward probability, transition probability, observation likelihood and backward probability:

$$\text{not - quite - } \xi_t(i, j) = \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j) \quad (17)$$

- To get ξ from not-quite- ξ , we have to divide by $P(O | \lambda)$, following the laws of probability:

$$P(X | Y, Z) = \frac{P(X, Y | Z)}{P(Y | Z)} \quad (18)$$

Training: The Forward-Backward Algorithm (VI)

- Define ξ_t as the probability of being in state i at time t and state j at time $t + 1$, given the observation sequence O and the HMM λ
- First compute something close to ξ , but including the probability of the observation:

$$\text{not - quite - } \xi_t(i, j) = P(q_t = i, q_{t+1} = j, O | \lambda) \quad (16)$$

- Use forward probability, transition probability, observation likelihood and backward probability:

$$\text{not - quite - } \xi_t(i, j) = \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j) \quad (17)$$

- To get ξ from not-quite- ξ , we have to divide by $P(O | \lambda)$, following the laws of probability:

$$P(X | Y, Z) = \frac{P(X, Y | Z)}{P(Y | Z)} \quad (18)$$

Training: The Forward-Backward Algorithm (VII)

- $P(O|\lambda)$ is simply the forward (or backward) probability of the whole utterance:

$$P(O|\lambda) = a_T(N) = \beta_T(1) = \sum_{j=1}^N \alpha_t(j) \beta_t(j) \quad (19)$$

- Therefore the final equation for ξ is:

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\alpha_T(N)} \quad (20)$$

- and the **expected number of transitions** from state i to j is the sum over all t of ξ

Training: The Forward-Backward Algorithm (VII)

- $P(O|\lambda)$ is simply the forward (or backward) probability of the whole utterance:

$$P(O|\lambda) = a_T(N) = \beta_T(1) = \sum_{j=1}^N \alpha_t(j) \beta_t(j) \quad (19)$$

- Therefore the final equation for ξ is:

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\alpha_T(N)} \quad (20)$$

- and the **expected number of transitions** from state i to j is the sum over all t of ξ

Training: The Forward-Backward Algorithm (VII)

- To get the **total expected number of transitions** from state i we sum over all transitions out of state i :

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{j=1}^N \xi_t(i, j)} \quad (21)$$

- We also need to recompute the **observation probability** $\hat{b}_j(v_k)$ (probability of a given symbol v_k from observation vocabulary V , given state j):

$$\hat{b}_j(v_k) = \frac{\text{expected number of times in state } j \text{ and observing symbol } v_k}{\text{expected number of times in state } j} \quad (22)$$

Training: The Forward-Backward Algorithm (VII)

- To get the **total expected number of transitions** from state i we sum over all transitions out of state i :

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{j=1}^N \xi_t(i, j)} \quad (21)$$

- We also need to recompute the **observation probability** $\hat{b}_j(v_k)$ (probability of a given symbol v_k from observation vocabulary V , given state j):

$$\hat{b}_j(v_k) = \frac{\text{expected number of times in state } j \text{ and observing symbol } v_k}{\text{expected number of times in state } j} \quad (22)$$

Training: The Forward-Backward Algorithm (VIII)

- We need to know the probability of being in state j at time t :

$$\gamma_t(j) = P(q_t = j | O, \lambda) = \frac{P(q_t = j, O | \lambda)}{P(O | \lambda)} = \frac{\alpha_t(j)\beta_t(j)}{P(O | \lambda)} \quad (23)$$

- Now we can compute b :
 - Numerator: sum $\gamma_t(j)$ for all time steps t in which the observation o_t is the symbol v_k
 - Denominator: sum $\gamma_t(j)$ over all time steps t

$$\hat{b}_j(v_k) = \frac{\sum_{t=1}^T \text{s.t. } O_t=v_k \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)} \quad (24)$$

Training: The Forward-Backward Algorithm (VIII)

- We need to know the probability of being in state j at time t :

$$\gamma_t(j) = P(q_t = j | O, \lambda) = \frac{P(q_t = j, O | \lambda)}{P(O | \lambda)} = \frac{\alpha_t(j)\beta_t(j)}{P(O | \lambda)} \quad (23)$$

- Now we can compute b :
 - Numerator: sum $\gamma_t(j)$ for all time steps t in which the observation o_t is the symbol v_k
 - Denominator: sum $\gamma_t(j)$ over all time steps t

$$\hat{b}_j(v_k) = \frac{\sum_{t=1}^T \text{s.t. } O_t=v_k \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)} \quad (24)$$

Training: The Forward-Backward Algorithm (IX)

- Now we can **re-estimate** A and B from O assuming that we have a previous estimate of A and B
- Forward-Backward Algorithm as a special case of EM algorithm:
 - ① E-step (expectation step):

Compute the expected state occupancy count γ and the expected state transition count ξ , from earlier A and B probabilities
 - ② M-step (maximisation step):

Use γ and ξ to recompute new A and B probabilities

Training: The Forward-Backward Algorithm (X)

function FORWARD-BACKWARD(*observations* of len T , *output vocabulary* V , *hidden state set* Q) **returns** $HMM=(A,B)$

initialize A and B

iterate until convergence

E-step

$$\gamma_t(j) = \frac{\alpha_t(j)\beta_t(j)}{P(O|\lambda)} \quad \forall t \text{ and } j$$

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\alpha_T(N)} \quad \forall t, i, \text{ and } j$$

M-step

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{j=1}^N \xi_t(i, j)}$$

$$\hat{b}_j(v_k) = \frac{\sum_{t=1s.t. O_t=v_k}^T \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)}$$

return A, B

Summary

Problem	Algorithm	Complexity
Evaluation: Calculating $P(q_t = q_i O_1, O_2 \dots O_t)$	Forward	$O(TN^2)$
Decoding: Computing $Q^* = \operatorname{argmax}_Q P(Q O)$	Viterbi	$O(TN^2)$
Learning: Computing $\lambda^* = \operatorname{argmax}_\lambda P(O \lambda)$	Forward-Backward (EM)	$O(TN^2)$

- HMMs have proved to be of great value in analysing real systems; their usual drawback is the over-simplification associated with the Markov assumption - that a state is dependent only on predecessors, and that this dependence is time independent.

References

- Jurafsky, D. and J. H. Martin. Speech and Language Processing. An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition. 2007.
- Online Tutorial: Hidden Markov Models.
(http://www.comp.leeds.ac.uk/roger/HiddenMarkovModels/html_dev/main.html)